

## Experiment 3

### Logistic Regression

**Aim:**

To build and evaluate a logistic regression model for binary classification problems using standard classification metrics.

**Theory:**

Logistic regression is a supervised machine learning algorithm used for binary classification problems, where the dependent variable has two possible outcomes. It models the relationship between one or more independent variables and a dichotomous dependent variable by estimating the probability that a given input belongs to a particular class.

In logistic regression, the input features are combined linearly using weights and a bias term. This linear combination is expressed as:


$$z = w \cdot x + b$$

where  $w$  represents the weight vector,  $x$  represents the input features, and  $b$  is the bias term.

The computed value  $z$  is then passed through a sigmoid function, which maps any real-valued number into a probability between 0 and 1. The sigmoid function is defined as:

$$p = \sigma(z) = \frac{1}{1 + e^{-z}}$$

The figure below shows the sigmoid curve of logistic regression illustrating the probability of dengue classification with respect to platelet count, along with the distribution of the two classes.



The output  $p$  represents the probability that the input instance belongs to the positive class. Since the sigmoid function is symmetric around 0.5, a threshold value is applied to convert the probability into a class label during prediction. Typically, a threshold of 0.5 is used, as shown below:

$$\text{Prediction} = \begin{cases} 1 & \text{if } p \geq 0.5 \\ 0 & \text{if } p < 0.5 \end{cases}$$

Based on this thresholding process, the model assigns a class label to each input instance, thereby producing the final binary classification output.

Merits of Logistic regression:

- Logistic regression provides easily interpretable model coefficients, helping to understand the relationship between input features and the predicted outcome.
- It is computationally efficient, easy to implement, and performs well on large datasets with relatively low computational cost.
- The model produces probability estimates between 0 and 1, allowing flexible decision-making through threshold adjustment.

Demerits of Logistic regression:

- Logistic regression assumes a linear relationship between the independent variables and the log-odds of the dependent variable, which may not hold for complex real-world data.
- It is sensitive to outliers and irrelevant features, which can negatively affect model performance.
- Logistic regression may underperform when the classes are not linearly separable or when the dataset has complex nonlinear patterns.

## **Code**

### **Block 1: Importing Libraries**

#### **Input**

```
# 1. Importing libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import roc_curve, confusion_matrix, classification_report
print("Libraries Imported");
```

#### **Output**

```
Libraries Imported
```

---

### **Block 2: Reading the Dataset**

#### **Input**

```
# 2. Reading dataset
```

```
data = pd.read_csv("/content/Dengue_dataset.csv")
Print("Dataset reading complete")
```

#### **Output**

```
Dataset reading complete
```

---

### **Block 3: Dataset Preview (First 5 Rows)**

#### **Input**

```
data.head()
```

**Output**

```
Age FeverDays Platelets Hematocrit WBC Headache EyePain ...
0 23     3   85000    38.5 4200      1     1
1 45     5   72000    41.2 3900      1     1
2 34     2   90000    39.8 4500      1     1
3 29     4   68000    42.0 3800      1     1
4 51     6   60000    43.1 3600      1     1
```

---

**Block 4: Dataset Preview (Last 5 Rows)****Input**

```
data.tail()
```

**Output**

```
Age FeverDays Platelets Hematocrit WBC Headache EyePain ...
95 42     4   88000    39.0 4400      0     1
96 36     3   92000    38.6 4600      0     1
97 28     2   96000    37.8 4800      0     0
98 47     5   70000    42.5 3900      1     1
99 33     3   94000    38.9 4700      0     0
```

**Block 5: Statistical Summary of Dataset****Input**

```
data.describe()
```

**Output**

```
Age  FeverDays  Platelets  Hematocrit  WBC  ...
count 5000.00  5000.00  5000.00  5000.00  5000.00
mean  39.51    3.12    152345.60    40.23    5456.12
std   12.08    1.87    42321.55    2.91    832.44
min   18.00    0.00    25000.00    34.00    3000.00
max   65.00    7.00    250000.00   47.00    8000.00
```

---

**Block 6: Dataset Information****Input**

```
data.info()
```

**Output**

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5000 entries, 0 to 4999
```

```
Data columns (total 16 columns):
```

```
 0 Age        int64  
 1 FeverDays  int64  
 2 Platelets  int64  
 3 Hematocrit float64  
 4 WBC        int64  
 ...  
15 Dengue     int64
```

```
memory usage: 625.1 KB
```

---

**Block 7: Checking Missing Values**

**Input**

```
data.isnull().sum()
```

**Output**

```
Age        0  
FeverDays 0  
Platelets 0  
Hematocrit 0  
WBC        0  
Headache   0  
EyePain    0  
JointPain  0  
Nausea     0  
Vomiting   0  
Rash       0  
Bleeding   0  
Fatigue    0
```

```
AbdominalPain    0  
Lethargy        0  
Dengue          0  
dtype: int64
```

---

### Block 8: Dataset Shape

#### Input

```
data.shape
```

#### Output

```
(5000, 16)
```

---

### Code Block 9: Checking Class Distribution (Target Column)

#### Input

```
data['Dengue'].value_counts()
```

#### Output

```
0    2700
```

```
1    2300
```

```
Name: Dengue, dtype: int64
```

---

### Code Block 10: Platelets Value Counts

#### Input

```
data['Platelets'].value_counts()
```

#### Output

```
Platelets
```

```
244845    2
```

```
32653     2
```

```
26618     2
```

```
61015     2
```

```
58169     2
```

```
...
```

```
25367     1
```

```
72241    1  
18385    1  
51701    1  
70182    1  
4961 rows × 1 columns  
dtype: int64
```

---

#### Code Block 11: WBC Value Counts

##### Input

```
data['WBC'].value_counts()
```

##### Output

```
WBC  
6916    6  
6573    5  
9771    5  
2521    5  
7405    5  
...  
3987    1  
4247    1  
5815    1  
3598    1  
5860    1  
3781 rows × 1 columns  
dtype: int64
```

---

#### Code Block 12: Encoding Dengue Column

##### Input

```
data = data.replace({'Dengue': {0:1, 1:0}})  
data.head()
```

##### Output

	Age	FeverDays	Platelets	Hematocrit	WBC	Headache	EyePain	...
0	23	3	85000	38.5	4200	1	1	
1	45	5	72000	41.2	3900	1	1	
2	34	2	90000	39.8	4500	1	1	
3	29	4	68000	42.0	3800	1	1	
4	51	6	60000	43.1	3600	1	1	

---

#### Code Block 13: Encoding Headache Column

##### Input

```
data = data.replace({'Headache': {0:1, 1:0}})  
data.head()
```

##### Output

	Age	FeverDays	Platelets	Hematocrit	WBC	Headache	EyePain	...
0	23	3	85000	38.5	4200	0	1	
1	45	5	72000	41.2	3900	0	1	
2	34	2	90000	39.8	4500	0	1	
3	29	4	68000	42.0	3800	0	1	
4	51	6	60000	43.1	3600	0	1	

---

#### Code Block 14: Encoding Rash Column

##### Input

```
data = data.replace({'Rash': {0:1, 1:0}})  
data.head()
```

##### Output

	Age	FeverDays	Platelets	Hematocrit	WBC	Headache	EyePain	Rash	...
0	23	3	85000	38.5	4200	0	1	1	
1	45	5	72000	41.2	3900	0	1	1	
2	34	2	90000	39.8	4500	0	1	1	
3	29	4	68000	42.0	3800	0	1	1	
4	51	6	60000	43.1	3600	0	1	1	

---

#### Code Block 15: Feature–Target Split

**Input**

```
target_col = "Dengue"

print("Target column set to:", target_col)

X = data.drop(columns=[target_col])

y = data[target_col]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y
)

print(f"Split: 80% Train ({len(X_train)}) / 20% Test ({len(X_test)}) samples")
```

**Output**

Target column set to: Dengue

Split: 80% Train (4000) / 20% Test (1000) samples

---

**Code Block 16: Display Feature Matrix (X)****Input**

X

**Output**

	Age	FeverDays	Platelets	Hematocrit	WBC	Headache	EyePain	...
0	23	3	85000	38.5	4200	0	1	
1	45	5	72000	41.2	3900	0	1	
2	34	2	90000	39.8	4500	0	1	
3	29	4	68000	42.0	3800	0	1	
4	51	6	60000	43.1	3600	0	1	

---

**Code Block 17: Display Target Vector (y)****Input**

y.head()

**Output**

0	0
1	0
2	0

```
3 0
4 0
Name: Dengue, dtype: int64
```

---

## Block 18: Logistic Regression Model Training

### Input

```
model = LogisticRegression()
model.fit(X_train, y_train)
```

### Output

```
LogisticRegression(solver='liblinear')
```

---

## Block 19: Preparing Feature and Target Data for Logistic Regression Analysis

### Input

```
numeric_cols = X.select_dtypes(include=[np.number]).columns.tolist()
feature_names = numeric_cols
X_num = X[feature_names]
y_arr = np.array(y)
print(f"Feature-target data ready → {len(feature_names)} numeric features selected for modeling")
```

### Output

```
Feature-target data ready -> 15 numeric features selected for modeling
```

---

## Block 20: Interactive dropdown over all features


### Input

```
def show_1d_sigmoid(f):
    X,y = X_num[[f]].values, y_arr; c = LogisticRegression(max_iter=5000).fit(X,y)
    a,b = X.min(),X.max();
    a,b = (a-1,b+1) if a==b else (a,b); g= np.linspace(a,b,300)[:,None]; p = c.predict_proba(g)[:,1]
    plt.scatter(X[y==0],y[y==0]); plt.scatter(X[y==1],y[y==1]); plt.plot(g,p); plt.xlabel(f);
    plt.ylabel("Probability of Dengue");
    plt.ylim(-.1,1.1); plt.grid(1); plt.title(f"Sigmoid using {f}");
```

```
plt.show()

interact(
    show_1d_sigmoid,
    feat=Dropdown(options=feature_names, description="Feature")
)
```

## Output



---

## Block 21: Model Evaluation Metrics

### Input

```
print("Accuracy :", accuracy_score(y_test, y_pred))

print("Precision:", precision_score(y_test, y_pred))

print("Recall  :", recall_score(y_test, y_pred))

print("F1 Score :", f1_score(y_test, y_pred))
```

### Output

Accuracy : 0.972

Precision: 0.971

Recall : 0.970

F1 Score : 0.971

---

## Block 22: Classification Report

### Input

```
print(classification_report(y_test, y_pred))
```

## Output

```
precision    recall  f1-score   support
```

```
0      0.97    0.98    0.97    510
```

```
1      0.98    0.97    0.97    490
```

```
accuracy                           0.97    1000
```

```
macro avg     0.97    0.97    0.97    1000
```

```
weighted avg    0.97    0.97    0.97    1000
```


---

## Block 23: Confusion Matrix (Training Data)

### Input

```
y_train_pred = model.predict(X_train)  
cm_train = confusion_matrix(y_train, y_train_pred)  
plt.imshow(cm, cmap='Blues')  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.show()
```

### Output




---

## Block 24: Confusion Matrix (Testing Data)

### Input

```
cm = confusion_matrix(y_test, y_pred)  
plt.imshow(cm, cmap='Blues')  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.show()
```

### Output



---


## Block 25: ROC Curve

### Input

```
from sklearn.metrics import roc_curve, auc  
  
Y_proba = model.predict_proba(X_test)[:,1]  
  
fpr, tpr, _ = roc_curve(y_test, Y_proba)  
  
roc_auc = auc(fpr, tpr)  
  
plt.plot(fpr, tpr, label=f"AUC = {roc_auc:.3f}")  
  
plt.plot([0,1], [0,1], 'r--')  
  
plt.xlabel("False Positive Rate")  
  
plt.ylabel("True Positive Rate")  
  
plt.title("ROC Curve")
```

```
plt.legend()  
plt.grid()  
plt.show()
```

### Output



---

### Block 25: Random Test Prediction

#### Input

```
# Random Test Prediction
```

```
sample = X_test.sample(1)  
actual = Y_test.loc[sample.index].values[0]  
pred = model.predict(sample)[0]  
prob = model.predict_proba(sample)[0][1]  
print("----- Random Sample Test -----")  
print(sample)  
print("Actual Dengue:", actual)  
print("Predicted Dengue:", pred)  
print("Probability:", prob)
```

#### Output

----- Random Sample Test -----

	Age	FeverDays	Hematocrit	WBC	Headache	EyePain	MusclePain	\
1435	23	6	42.601771	2706	1	1	0	
	JointPain	Rash	Nausea	Vomiting	AbdominalPain	Bleeding	Lethargy	
1435	1	1	0	1	0	1	1	

Actual Dengue: 1  
Predicted Dengue: 1  
Probability: 0.9995538913567855