

Logistic Regression - Practical Implementation

Virtual Labs - Dayalbagh Educational Institute

Step 1: Importing Libraries

Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, roc_curve, confusion_matrix, classification_report
)
from ipywidgets import interact, Dropdown
print("Libraries Imported");
```

Output:

Libraries Imported

Step 2: Loading Dataset

Read and store the Dengue dataset

```
data = pd.read_csv('Dengue_dataset.csv')
print("Dataset reading complete")
```

Output:

Dataset reading complete

Step 3: Data Analysis

Display the first 5 rows of the dataset

```
data.head()
```

Output:

index	Age	Fever Days	Platelets	Hematocrit	WBC	Headache	EyePain	Muscle Pain	Joint Pain	Rash	Nausea	Vomiting	Abdominal Pain	Bleeding	Lethargy	Dengue
0	61	6	56777	49.49572857	3205	0	1	0	0	0	1	1	0	0	0	1
1	24	7	61250	42.41093608	3734	1	0	0	0	1	0	0	1	1	0	1
2	70	7	88034	49.66143051	3057	1	0	1	0	0	1	0	0	0	0	1
3	30	6	55130	50.20159282	2711	1	0	0	1	0	0	1	1	0	0	1
4	33	3	64346	43.46980401	4884	1	1	0	1	1	1	1	0	1	0	1

Display the last 5 rows of the dataset

```
data.tail()
```

Output:

index	Age	Fever_Days	Platelets	Hematocrit	WBC	Headache	EyePain	Muscle_Pain	Joint_Pain	Rash	Nausea	Vomiting	Abdominal_Pain	Bleeding	Lethargy	Dengue
4995	59	2	123791	47.18861033	423	1	0	0	0	0	0	0	0	0	0	0
4996	11	4	111575	41.84560446	598	1	0	1	1	1	1	1	0	0	0	0
4997	19	2	128198	40.2842782	477	1	1	1	0	0	0	0	0	0	1	0
4998	11	0	141619	47.3939162	462	1	1	0	0	1	1	1	1	0	1	0
4999	69	2	90619	46.47148497	598	1	0	0	0	1	0	1	0	0	0	0

Show statistical summary of numerical columns

```
data.describe()
```

Output:

	Age	Fever_Days	Platelets	Hematocrit	WBC	Headache	Eye_Pain	Muscle_Pain	Joint_Pain	Rash	Nausea	Vomiting	Abdominal_Pain	Bleeding	Lethargy	Dengue
count	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0	5000.0
mean	41.85	3.41	141522.48	41.86	6156.58	0.517	0.492	0.508	0.500	0.496	0.501	0.507	0.496	0.226	0.502	0.460
std	18.80	2.05	87116.73	4.81	2594.06	0.500	0.500	0.500	0.500	0.500	0.500	0.500	0.500	0.418	0.500	0.498
min	10.0	0.0	15060.0	32.02	2002.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
25%	26.0	2.0	62365.5	38.43	3848.75	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
50%	42.0	3.0	130972.0	42.02	6169.0	1.0	0.0	1.0	0.5	0.0	1.0	1.0	0.0	0.0	1.0	0.0
75%	58.0	5.0	221014.25	44.89	8302.5	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0	1.0	1.0
max	74.0	7.0	299988.0	51.99	10993.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

Display dataset structure and data types

```
data.info()
```

Output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 16 columns):
```

Output:

#	Column	Non-Null Count	Dtype
0	Age	5000 non-null	int64
1	FeverDays	5000 non-null	int64
2	Platelets	5000 non-null	int64
3	Hematocrit	5000 non-null	float64
4	WBC	5000 non-null	int64
5	Headache	5000 non-null	int64
6	EyePain	5000 non-null	int64
7	MusclePain	5000 non-null	int64
8	JointPain	5000 non-null	int64
9	Rash	5000 non-null	int64
10	Nausea	5000 non-null	int64
11	Vomiting	5000 non-null	int64
12	AbdominalPain	5000 non-null	int64

13	Bleeding	5000 non-null	int64
14	Lethargy	5000 non-null	int64

... 1 more rows (total: 16 rows)

Output:

```
dtypes: float64(1), int64(15)
memory usage: 625.1 KB
```

Check the number of missing values in each column

```
data.isnull().sum()
```

Output:

Column	Missing Values
Age	0
FeverDays	0
Platelets	0
Hematocrit	0
WBC	0
Headache	0
EyePain	0
MusclePain	0
JointPain	0
Rash	0
Nausea	0
Vomiting	0
AbdominalPain	0
Bleeding	0
Lethargy	0

... 1 more rows (total: 16 rows)

Output:

```
dtype: int64
```

Show the number of rows and columns in the dataset

```
data.shape
```

Output:

```
(5000, 16)
```

Count occurrences of each class in the Dengue column

```
data['Dengue'].value_counts()
```

Output:

```
0    27001    2300Name: Dengue, dtype: int64
```

Count frequency of each Platelets value

```
data['Platelets'].value_counts()
```

Output:

```
Platelets
244845    2
32653     2
26618     2
61015     2
58169     2
...
25367     1
72241     1
18385     1
51701     1
70182     1
4961 rows x 1 columns
dtype: int64
```

Count frequency of each WBC value

```
data['WBC'].value_counts()
```

Output:

```
WBC
6916     6
6573     5
9771     5
2521     5
7405     5
...
3987     1
4247     1
5815     1
3598     1
5860     1
3781 rows x 1 columns
dtype: int64
```

Step 4: Data Preprocessing

Encode Dengue column by flipping values (0 -> 1, 1 -> 0)

```
data = data.replace({'Dengue':{0:1, 1:0}})
data.head()
```

Output:

Dataset Preview:

Output:

index	Age	FeverDays	Platelets	Hematocrit	WBC	Headache	EyePain	MusclePain	JointPain	Rash	Nausea	Vomiting	AbdominalPain	Bleeding	Lethargy	Dengue
0	61	6	56777	49.4957	3205	0	1	0	0	0	1	1	0	0	0	0
1	24	7	61250	42.4109	3738	1	0	0	0	1	0	0	1	1	0	0
2	70	7	88034	49.6614	3052	1	0	1	0	0	1	0	0	0	0	0

3	30	6	55130	50.2016	2713	1	0	0	1	0	0	1	1	0	0	0
4	33	3	64346	43.4698	4888	1	1	0	1	1	1	1	0	1	0	0

Encode Headache column by flipping values (0 -> 1, 1 -> 0)

```
data = data.replace({'Headache':{0:1, 1:0}})
data.head()
```

Output:

Dataset Preview:

Output:

index	Age	FeverDays	Platelets	Hematocrit	WBC	Headache	EyePain	MusclePain	JointPain	Rash	Nausea	Vomiting	AbdominalPain	Bleeding	Lethargy	Dengue
0	61	6	56777	49.4957	3205	1	1	0	0	0	1	1	0	0	0	0
1	24	7	61250	42.4109	3738	0	0	0	0	1	0	0	1	1	0	0
2	70	7	88034	49.6614	3052	0	0	1	0	0	1	0	0	0	0	0
3	30	6	55130	50.2016	2713	0	0	0	1	0	0	1	1	0	0	0
4	33	3	64346	43.4698	4888	0	1	0	1	1	1	1	0	1	0	0

Encode Rash column by flipping values (0 -> 1, 1 -> 0) and verify changes

```
data = data.replace({'Rash':{0:1, 1:0}})
print("Encoding complete. Checked head:")
data.head()
```

Output:

Encoding complete. Checked head:
Dataset Preview:

Output:

index	Age	FeverDays	Platelets	Hematocrit	WBC	Headache	EyePain	MusclePain	JointPain	Rash	Nausea	Vomiting	AbdominalPain	Bleeding	Lethargy	Dengue
0	61	6	56777	49.4957	3205	1	1	0	0	1	1	1	0	0	0	0
1	24	7	61250	42.4109	3738	0	0	0	0	0	0	0	1	1	0	0
2	70	7	88034	49.6614	3052	0	0	1	0	1	1	0	0	0	0	0
3	30	6	55130	50.2016	2713	0	0	0	1	1	0	1	1	0	0	0
4	33	3	64346	43.4698	4888	0	1	0	1	0	1	1	0	1	0	0

Step 5: Model Training

Set the target variable, separate features and labels, and split the dataset into 80% training and 20% testing data

```
target_col = "Dengue"
print("Target column set to:", target_col)
X = data.drop(columns=['Platelets', target_col])
y = data[target_col]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42,
                                                    stratify=y)
print(f"Split: 80% Train ({len(X_train)}) / 20% Test ({len(X_test)}) samples")
```

Output:

Target column set to: Dengue
Split: 80% Train (4000) / 20% Test (1000) samples

Display the feature matrix

X

Output:

Feature Matrix (X) Preview:

Output:

	Age	FeverDays	Platelets	Hematocrit	WBC	Headache	EyePain	MusclePain	JointPain	Rash	Nausea	Vomiting	AbdominalPain	Bleeding	Lethargy
0	61	6	56777	49.495729	3205	0	1	0	0	0	1	1	0	0	0
1	24	7	61250	42.410936	3738	1	0	0	0	1	0	0	1	1	0
2	70	7	88034	49.661431	3052	1	0	1	0	0	1	0	0	0	0
3	30	6	55130	50.201593	2713	1	0	0	1	0	0	1	1	0	0
4	33	3	64346	43.469804	4888	1	1	0	1	1	1	1	0	1	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
4995	59	2	123791	47.188610	4239	1	0	0	0	0	0	0	0	0	0
4996	11	4	111575	41.845604	5989	1	0	1	1	1	1	1	0	0	0
4997	19	2	128198	40.284278	4775	1	1	1	0	0	0	0	0	0	1
4998	11	0	141619	47.393916	4625	1	1	0	0	1	1	1	1	0	1
4999	69	2	90619	46.471485	5982	1	0	0	0	1	0	1	0	0	0

Output:

5000 rows x 15 columns

Display the first few values of the target variable

y.head()

Output:

Target Vector (y) Preview:

Output:

	Dengue
0	0
1	0
2	0
3	0
4	0

Output:

Name: Dengue, dtype: int64

Initialize and train a Logistic Regression model with L2 regularization on the training data

```
model = LogisticRegression(C=1.0, penalty='l2', solver='liblinear')
model.fit(X_train, Y_train)
print("Model trained successfully")
```

Output:

Model trained successfully

LogisticRegression



LogisticRegression(solver='liblinear')

Step 6: Model Evaluation

Preparing Feature and Target Data for Logistic Regression Analysis

```
numeric_cols = X.select_dtypes(include=[np.number]).columns.tolist()
feature_names = numeric_cols
X_num = X[feature_names]
y_arr = np.array(y)
print(f"Feature-target data ready -> {len(feature_names)} numeric features selected for modeling")
```

Output:

```
Feature-target data ready -> 14 numeric features selected for modeling
```

Fits a logistic regression model on a single feature and visualizes the resulting sigmoid probability curve against the

```
def show_1d_sigmoid(f):
    X,y = X_num[[f]].values, y_arr; c = LogisticRegression(max_iter=5000).fit(X,y)
    a,b = X.min(),X.max();
    a,b = (a-1,b+1) if a==b else (a,b); g= np.linspace(a,b,300)[:,:None]; p =
    c.predict_proba(g)[:,:1]
    plt.scatter(X[y==0],y[y==0]); plt.scatter(X[y==1],y[y==1]); plt.plot(g,p);
    plt.xlabel(f);
    plt.ylabel("Probability of Dengue");
    plt.ylim(-.1,1.1); plt.grid(1); plt.title(f"Sigmoid using {f}");
    plt.show()

print("Function defined: show_1d_sigmoid")
```

Output:


```
Function defined: show_1d_sigmoid
```

Interactive dropdown over all features

```
interact(
    show_1d_sigmoid,
    feat=Dropdown(options=feature_names, description="Feature")
)
```

Output:

Feature
 Age
 FeverDays
 Platelets
 Hematocrit
 WBC
 MusclePain
 JointPain
 Nausea
 Vomiting
 AbdominalPain
 Bleeding
 Lethargy
 Rash
 EyePain
 Headache

**# Training Performance**

```

Y_train_pred = model.predict(X_train)
print(" Training Performance ")
print(f"Accuracy : {accuracy_score(y_train, Y_train_pred):.4f}")
print(f"Precision: {precision_score(y_train, Y_train_pred):.4f}")
print(f"Recall   : {recall_score(y_train, Y_train_pred):.4f}")
print(f"F1 Score : {f1_score(y_train, Y_train_pred):.4f}")

```

Output:

Training Performance
 Accuracy : 0.9732
 Precision: 0.9712
 Recall : 0.9707
 F1 Score : 0.9709

Confusion Matrix for training data


```

cm_train = confusion_matrix(y_train, Y_train_pred)
sns.heatmap(cm_train, cmap="Blues")
plt.title("Confusion Matrix (Training Data)")
for i in range(len(cm_train)):
    for j in range(len(cm_train)):
        plt.text(j, i, cm_train[i][j], ha='center', va='center')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```

Output:

Confusion Matrix (Training Data)



Testing Performance

```

Y_test_pred = model.predict(X_test)
print("Testing Performance")
print(f"Accuracy : {accuracy_score(y_test, Y_test_pred):.4f}")
print(f"Precision: {precision_score(y_test, Y_test_pred):.4f}")
print(f"Recall : {recall_score(y_test, Y_test_pred):.4f}")
print(f"F1 Score : {f1_score(y_test, Y_test_pred):.4f}")

```

Output:


Testing Performance
Accuracy : 0.9740
Precision: 0.9657
Recall : 0.9783
F1 Score : 0.9719

Confusion Matrix for testing data

```
cm_test = confusion_matrix(y_test, Y_test_pred)
sns.heatmap(cm_test, annot=True, cmap="Blues", fmt='d')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix for Testing Data")
plt.show()
```

Output:

Confusion Matrix (Testing Data)

**# Calculates predicted probabilities to compute ROC curve metrics and the AUC score.**

```
Y_proba = model.predict_proba(X_test)[:,-1]
fpr, tpr, th = roc_curve(Y_test, Y_proba)
roc_auc = auc(fpr, tpr)
print(f"ROC metrics calculated. AUC = {roc_auc:.3f}")
```


Output:

ROC metrics calculated. AUC = 0.998

Plots the ROC curve including the AUC score and a random classifier baseline.

```
plt.plot(fpr, tpr, linewidth=3, label=f"AUC = {roc_auc:.3f}")
plt.plot([0,1], [0,1], 'r--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.grid()
plt.show()
```

Output:
ROC Curve



Random Test Prediction

```
test_samples=data.sample(5)
display(test_samples)
sample = X_test.sample(1)
actual = Y_test.loc[sample.index].values[0]
pred = model.predict(sample)[0]
prob = model.predict_proba(sample)[0][1]
print("----- Random Sample Test -----")
print(sample)
print("Actual Dengue:", actual)
print("Predicted Dengue:", pred)
print("Probability:", prob)
```

Output:
Select a row to see prediction:

Output:

Index	Age	FeverDays	Hematocrit	WBC	Headache	EyePain	MusclePain	JointPain	Rash	Nausea	Vomiting	AbdominalPain	Bleeding	Lethargy
-------	-----	-----------	------------	-----	----------	---------	------------	-----------	------	--------	----------	---------------	----------	----------


Interactive Features (Static View)

Interactive Feature: Sigmoid Curve Viewer

In the simulation, a dropdown allows you to select any feature and view its individual sigmoid curve. The available features are:

Platelets (default) | Age | FeverDays
 Hematocrit | WBC | MusclePain
 JointPain | Nausea | Vomiting
 AbdominalPain | Bleeding | Lethargy
 Rash | EyePain | Headache

Default view: Platelets sigmoid curve



Interactive Feature: Test Sample Predictions

Clicking a row in the test samples table shows the model prediction. Below are all 5 test samples with their prediction results:

Index	Age	Fever	Platelets	Hematocrit	WBC	Headache	EyePain	Muscle	Joint	Rash	Nausea	Vomit	Abdom	Bleed	Lethargy	Actual	Pred	Prob
1434	69	3	32794	44.36	4203	1	1	1	0	0	0	1	1	0	0	1	1	0.9995
1435	23	6	34881	42.60	2706	1	1	0	1	0	0	1	0	1	1	1	1	0.9782
1436	11	7	64703	48.04	3961	0	0	1	0	1	0	0	1	0	1	1	1	0.9810
1437	26	4	59354	51.65	4519	0	1	0	1	0	1	0	0	1	0	1	1	0.9755
1438	37	3	29390	42.11	4666	1	1	0	0	0	1	0	0	1	1	1	1	0.9991